

PHP & Kubernetes

Deploying right in your wheelhouse

Denys Bulakh



Denys Bulakh

Introduction

- Programmer
- Web Software Developer
- CTO at Regiondo
- Principal Engineer at Jochen Schweizer Mydays Group



JOCHEN
SCHWEIZER



mydays
ERLEBNISWERK



REGIONDO
Activity Booking Software

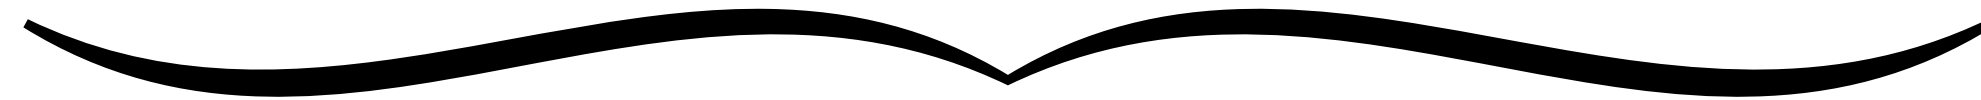
spontacts

What will we talk about?

Application deployment
Deployment automation
Application scaling
Local development
Logging



How to deploy web application?





1st challenge

**It runs permanently
on web server**



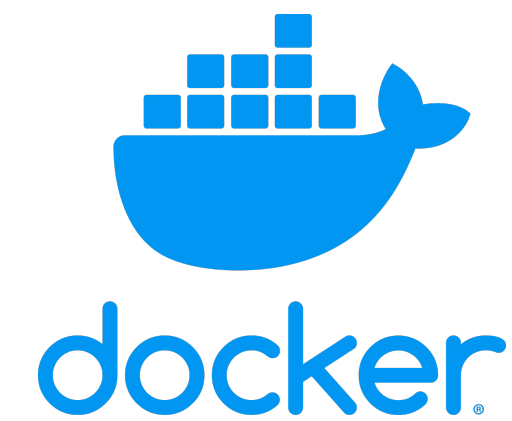
2nd challenge

Application versioning

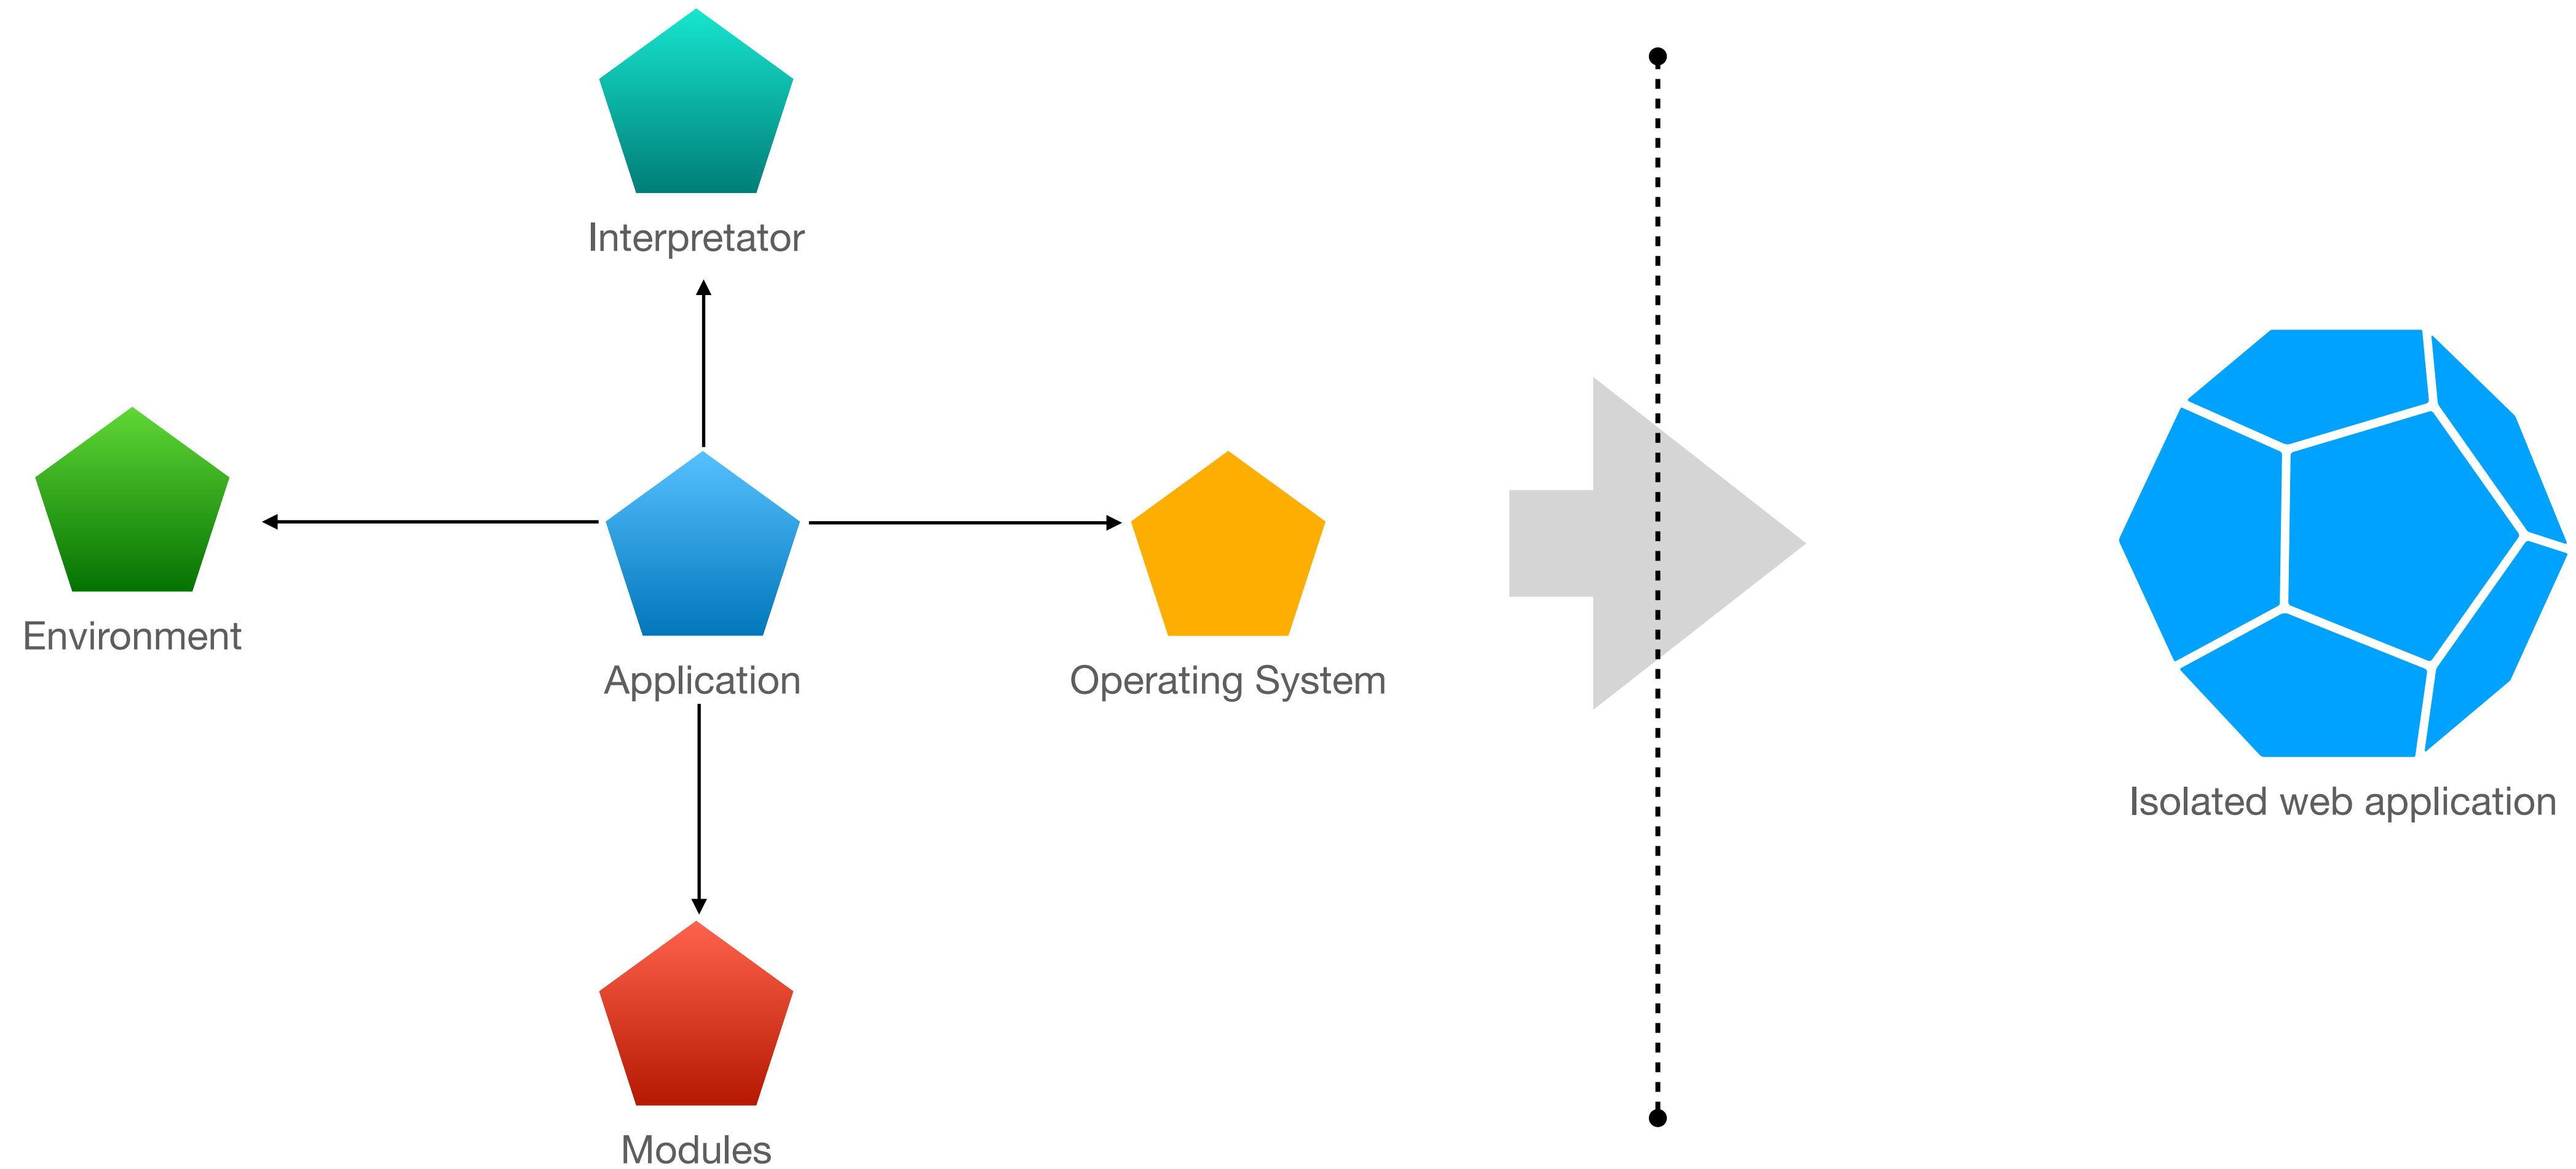


3rd challenge

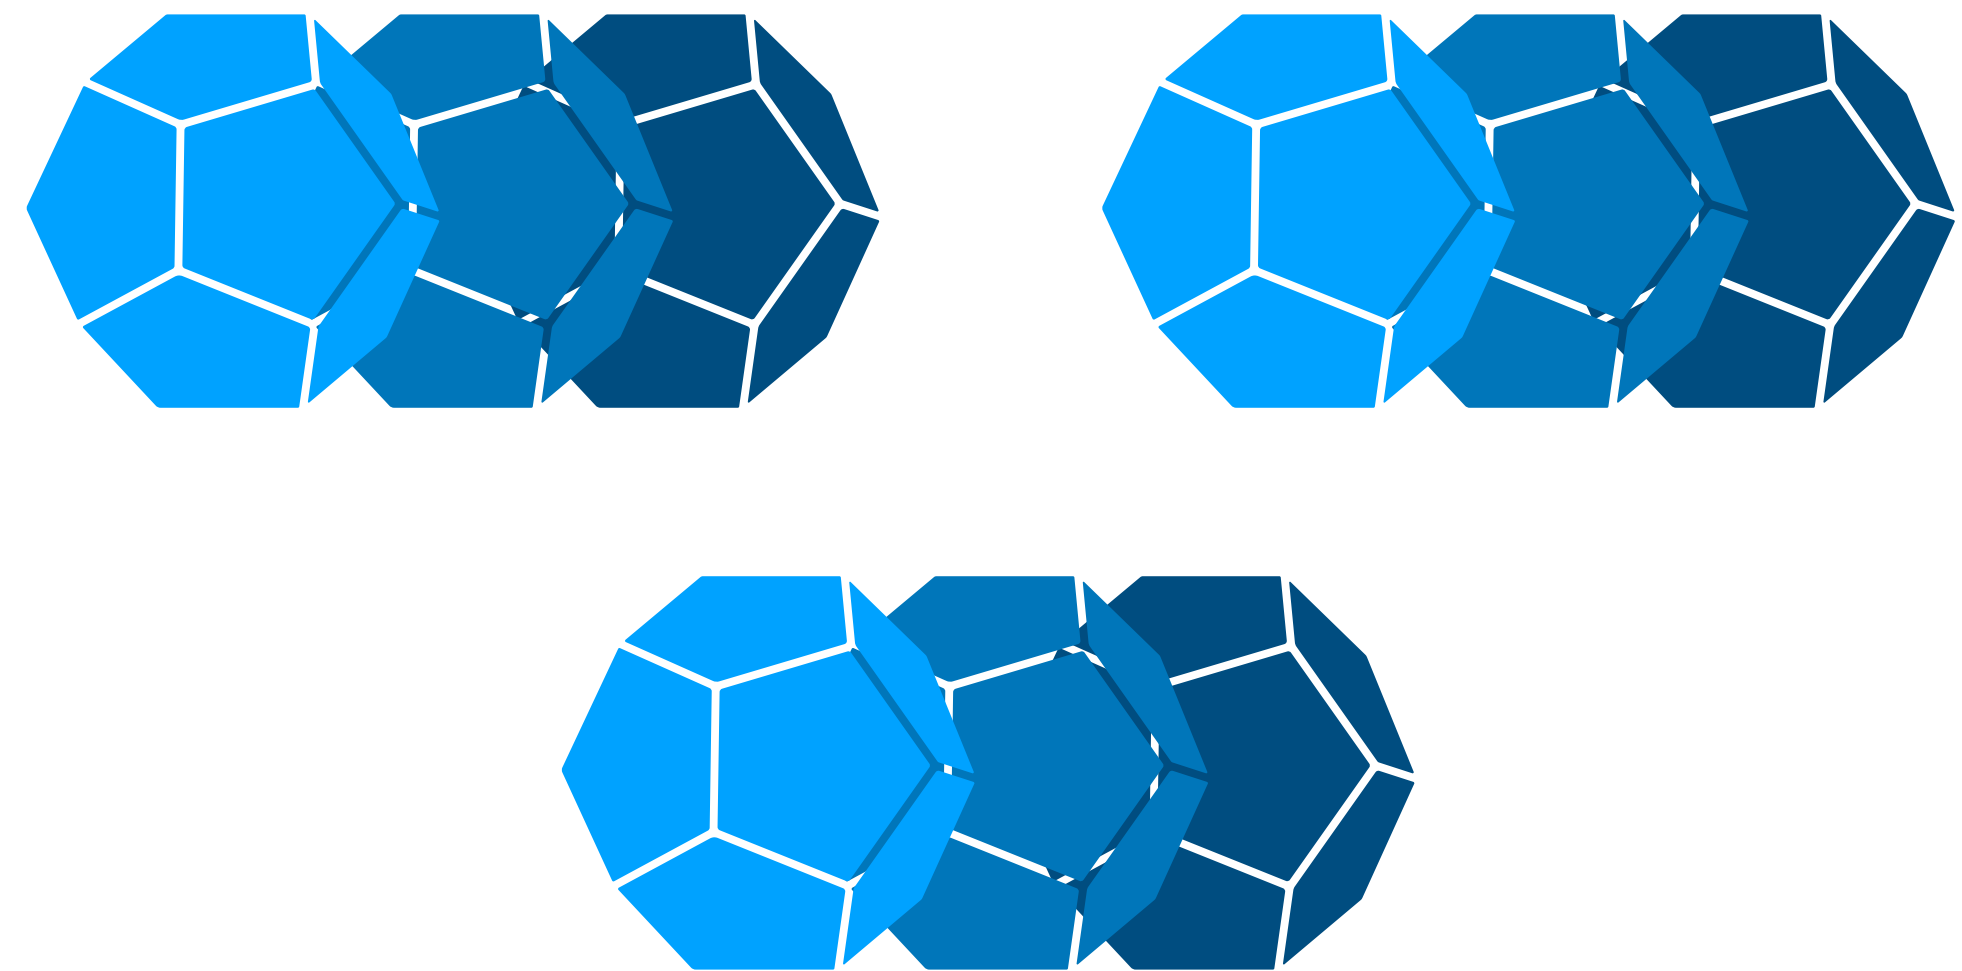
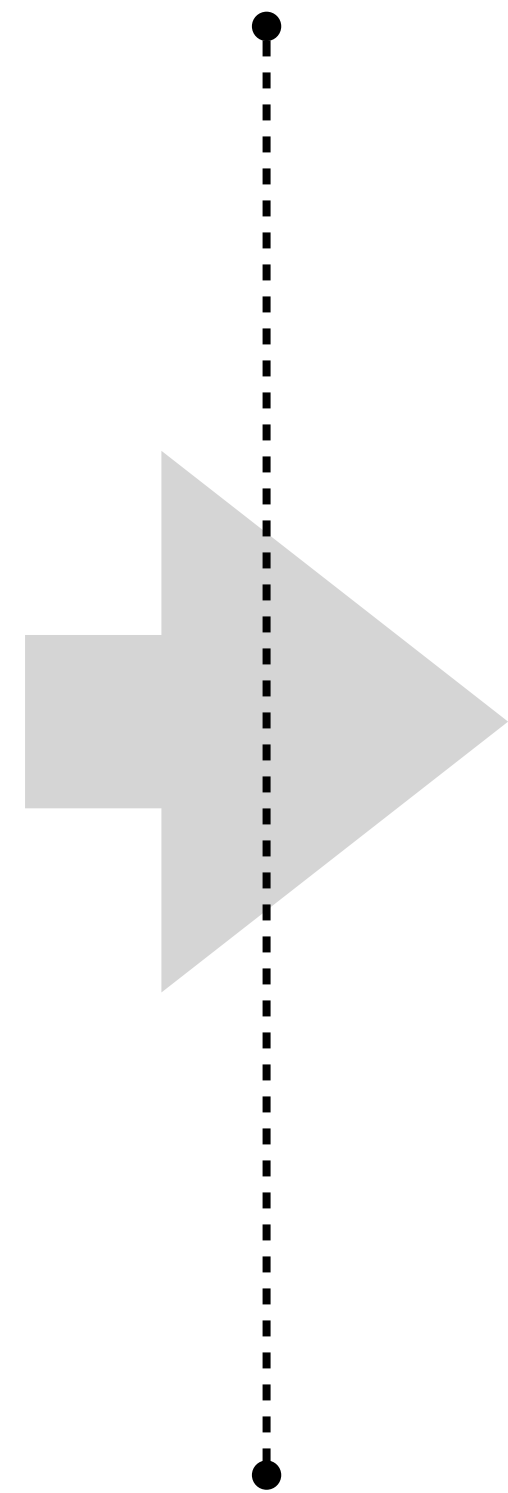
State management



**Application containerisation helps to
make stateless isolated web application**



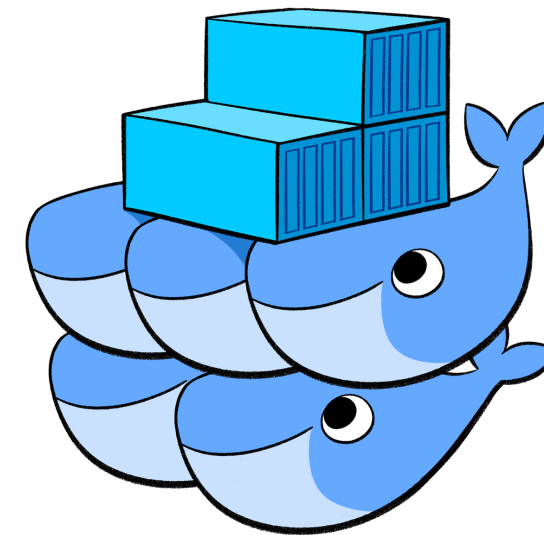
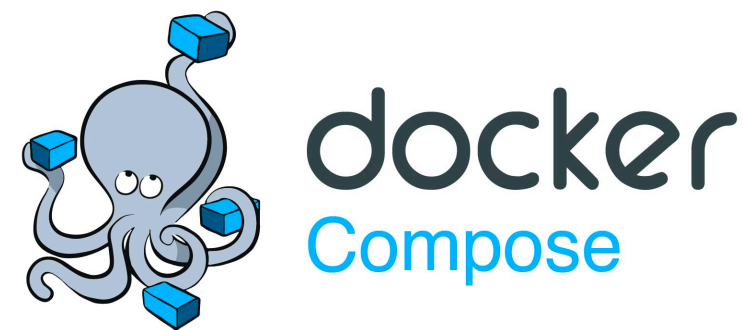
Instead of managing state and dependencies, we can include and isolate them in the application



We can run as many isolated web applications, as we need, without need to take care about state and dependencies



Orchestrators - tools to manage, scale, and maintain containerised applications



kubernetes



kubernetes

Features

- Automated rollouts and rollbacks
- Service discovery and load balancing
- Storage orchestration
- Secret and configuration management
- Horizontal scaling
- Self-healing
- IPv4/IPv6 dual-stack
- Managing resources for containers

<https://kubernetes.io/>

Kubernetes does not deploy source code and does not build your application

<https://kubernetes.io/>

Kubernetes does not provide application level services, such as databases, caches, middleware as build-in services

<https://kubernetes.io/>

**If an application can run in a container,
it should run great on Kubernetes**

<https://kubernetes.io/>

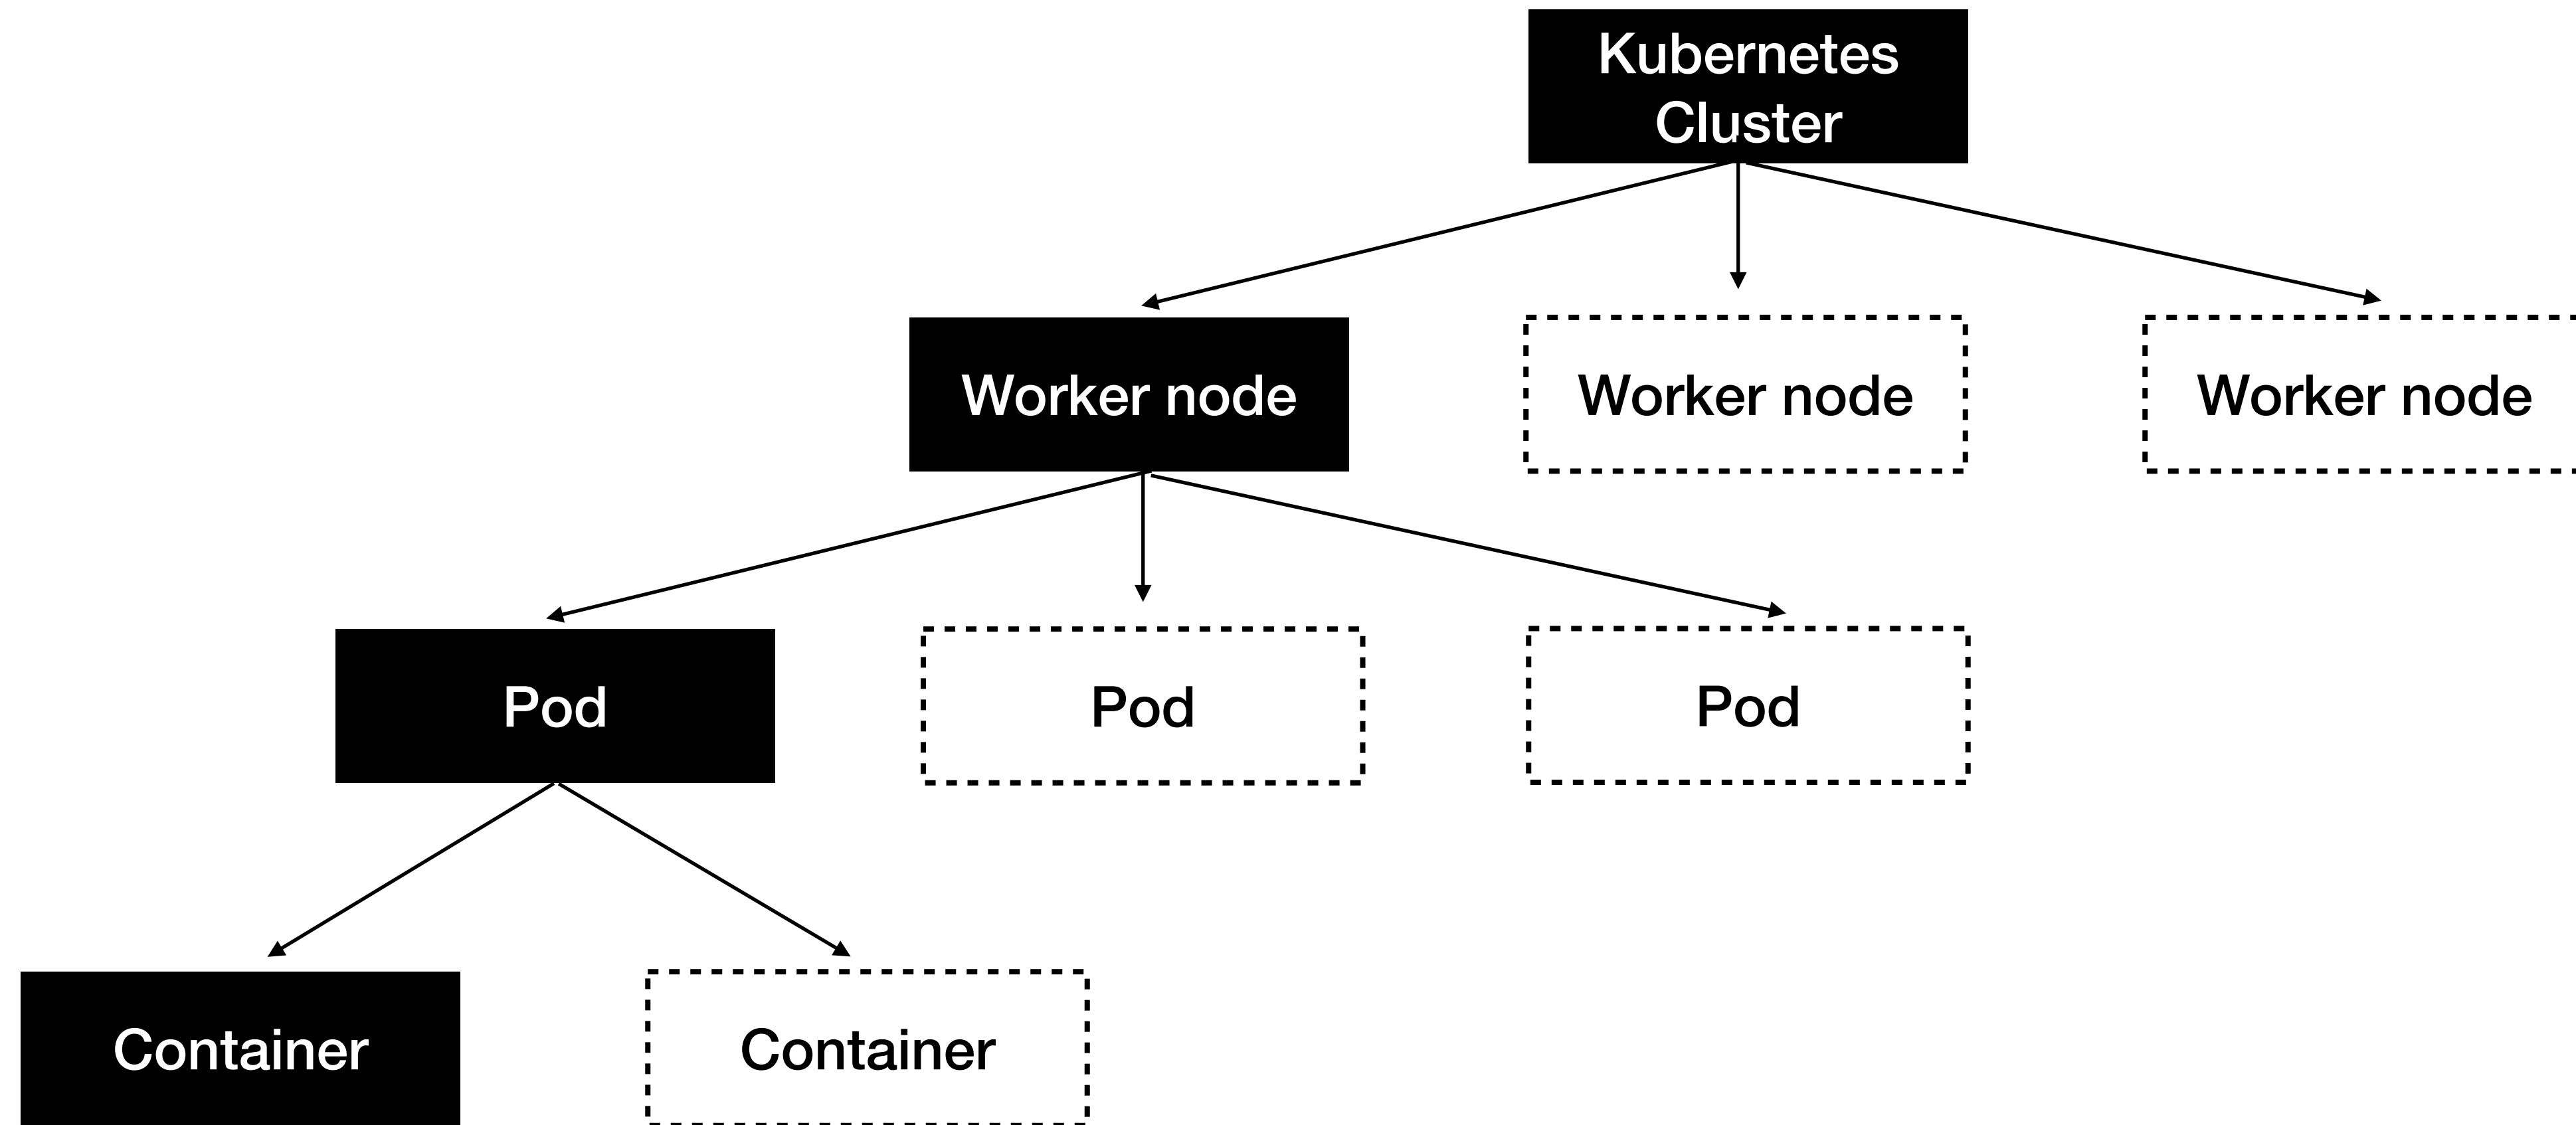
**PHP applications can be easily
containerised and, therefore, deployed
with Kubernetes**

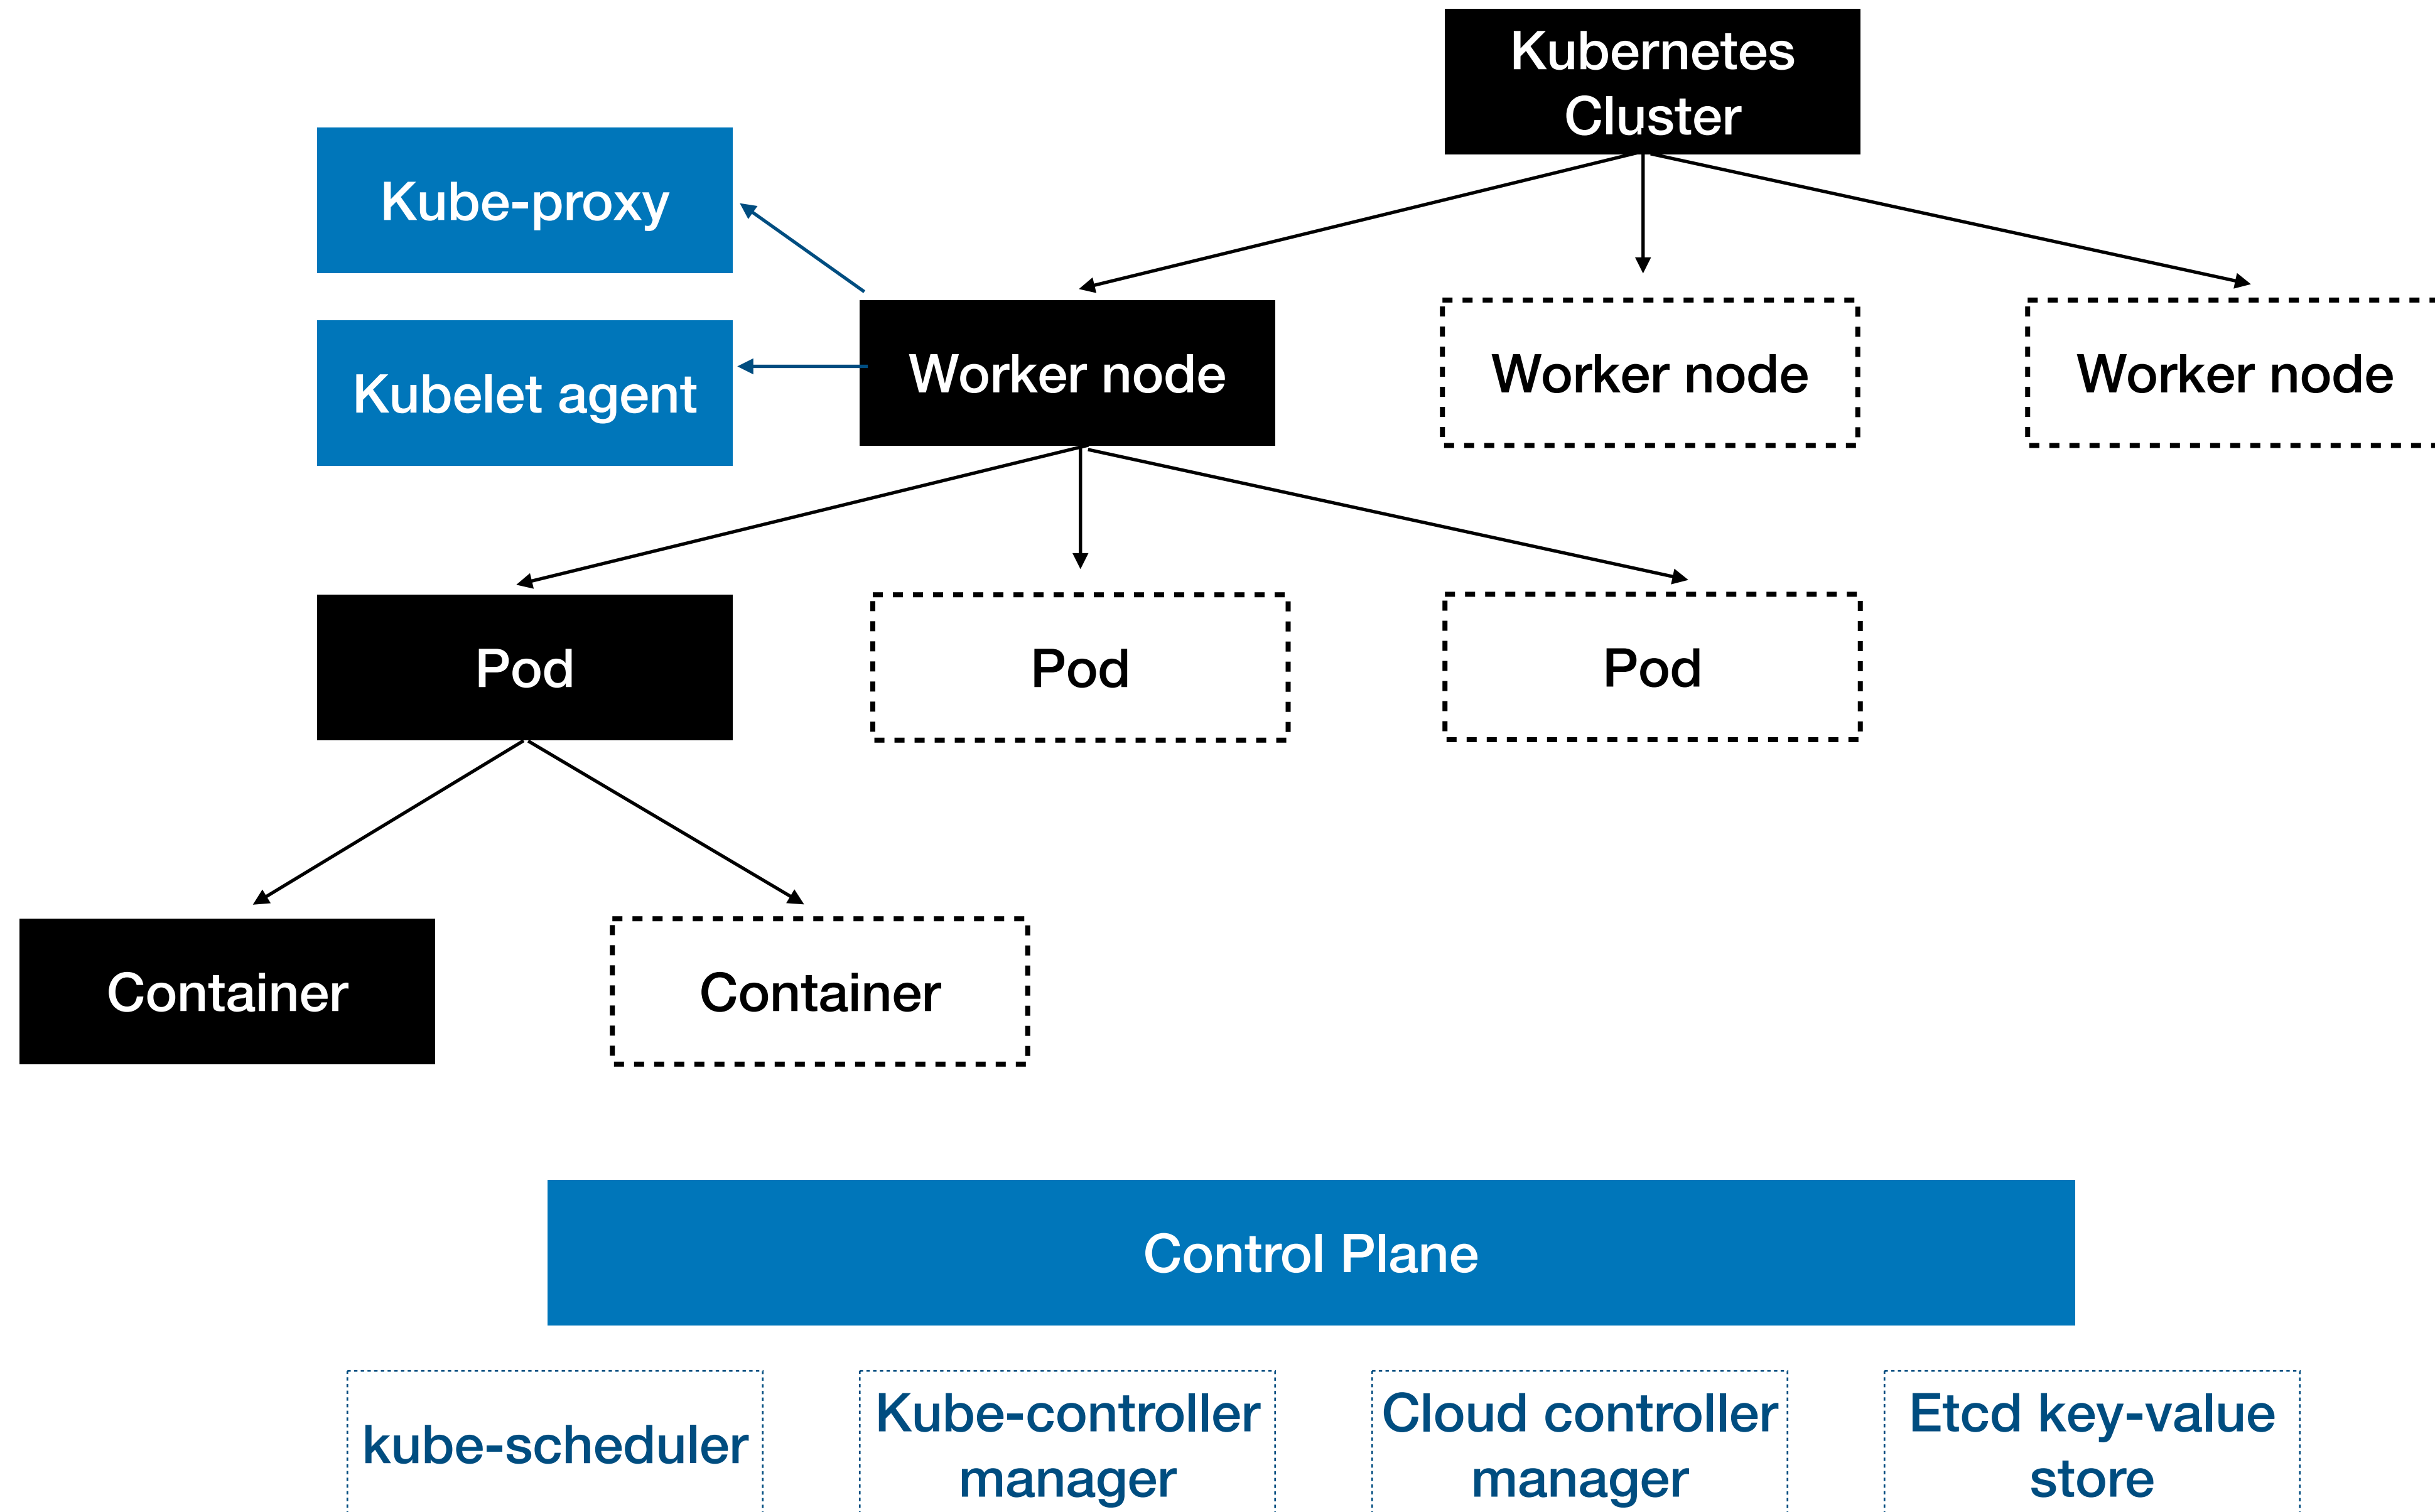


+



kubernetes





```
kubectl [command] [TYPE] [NAME] [flags]
```

Kubectl

Kubeadm

Rest API

API Server

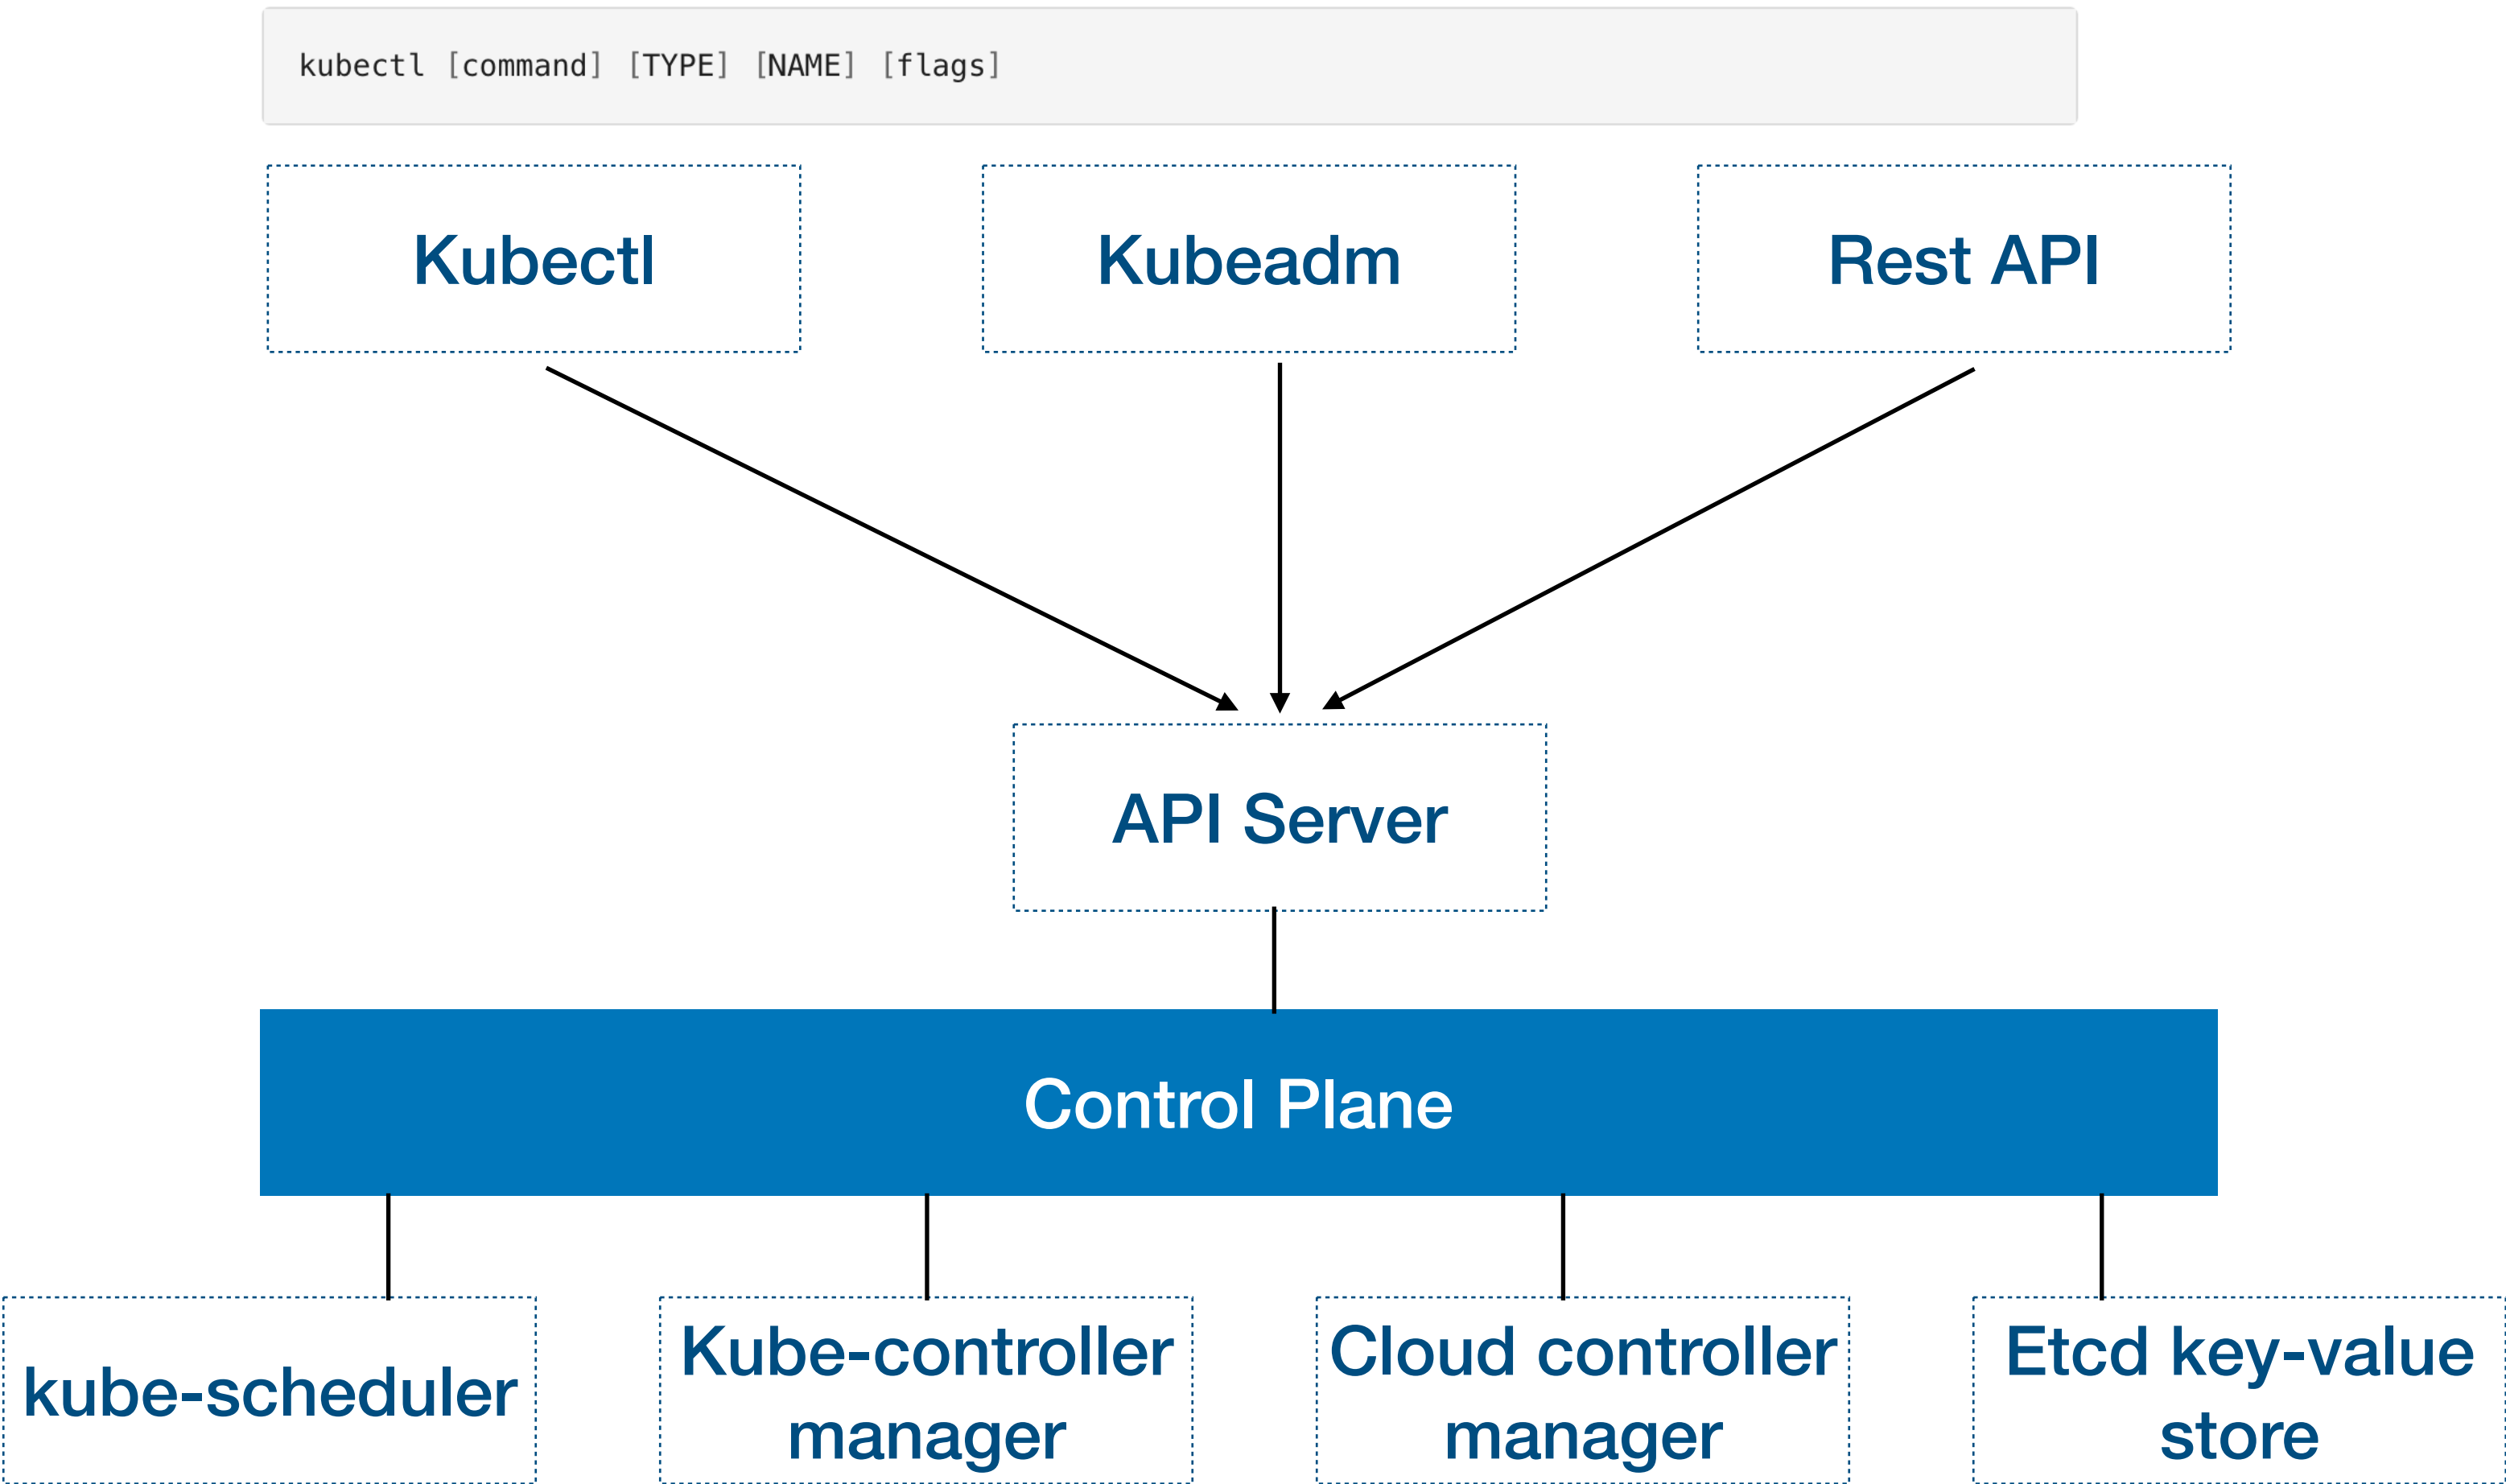
Control Plane

kube-scheduler

Kube-controller
manager

Cloud controller
manager

Etc
d key-value
store



Kubernetes is using objects, as persistent entities in the system, to represent the state of the cluster

Kubernetes Object

What and where is running?

Which resources are
available?

Restart, upgrade, fault-
tolerance policies

By creating an object, you're telling the Kubernetes system what do you want your cluster's workload to look like, this is your cluster's desired state

<https://kubernetes.io/>

Declare Kubernetes object's desired configuration

A

apiVersion: apps/v1 *# for versions before 1.9.0 use apps/v1beta2*

B

kind: Deployment

metadata:

name: nginx-deployment

spec:

selector:

matchLabels:

app: nginx

replicas: 2 *# tells deployment to run 2 pods matching the template*

template:

C metadata:

labels:

app: nginx

D spec:

containers:

– **name:** nginx

image: nginx:1.14.2

ports:

– **containerPort:** 80

- A. **version:** can be v1, v1beta and v2.
- B. **kind:** Pod, DaemonSet, Deployment, Service.
- C. **metadata:** name of the pod, namespace, labels, annotations.
- D. **spec:** desired state of the pod, including container image, replicas count, environment variables and volumes.

```
kubectl apply -f https://k8s.io/examples/application/deployment.yaml --record
```

<https://kubernetes.io/>

Other management techniques

Objects can be managed imperatively, by operating directly on live objects in cluster, using kubectl commands

```
kubectl create deployment nginx --image nginx
```

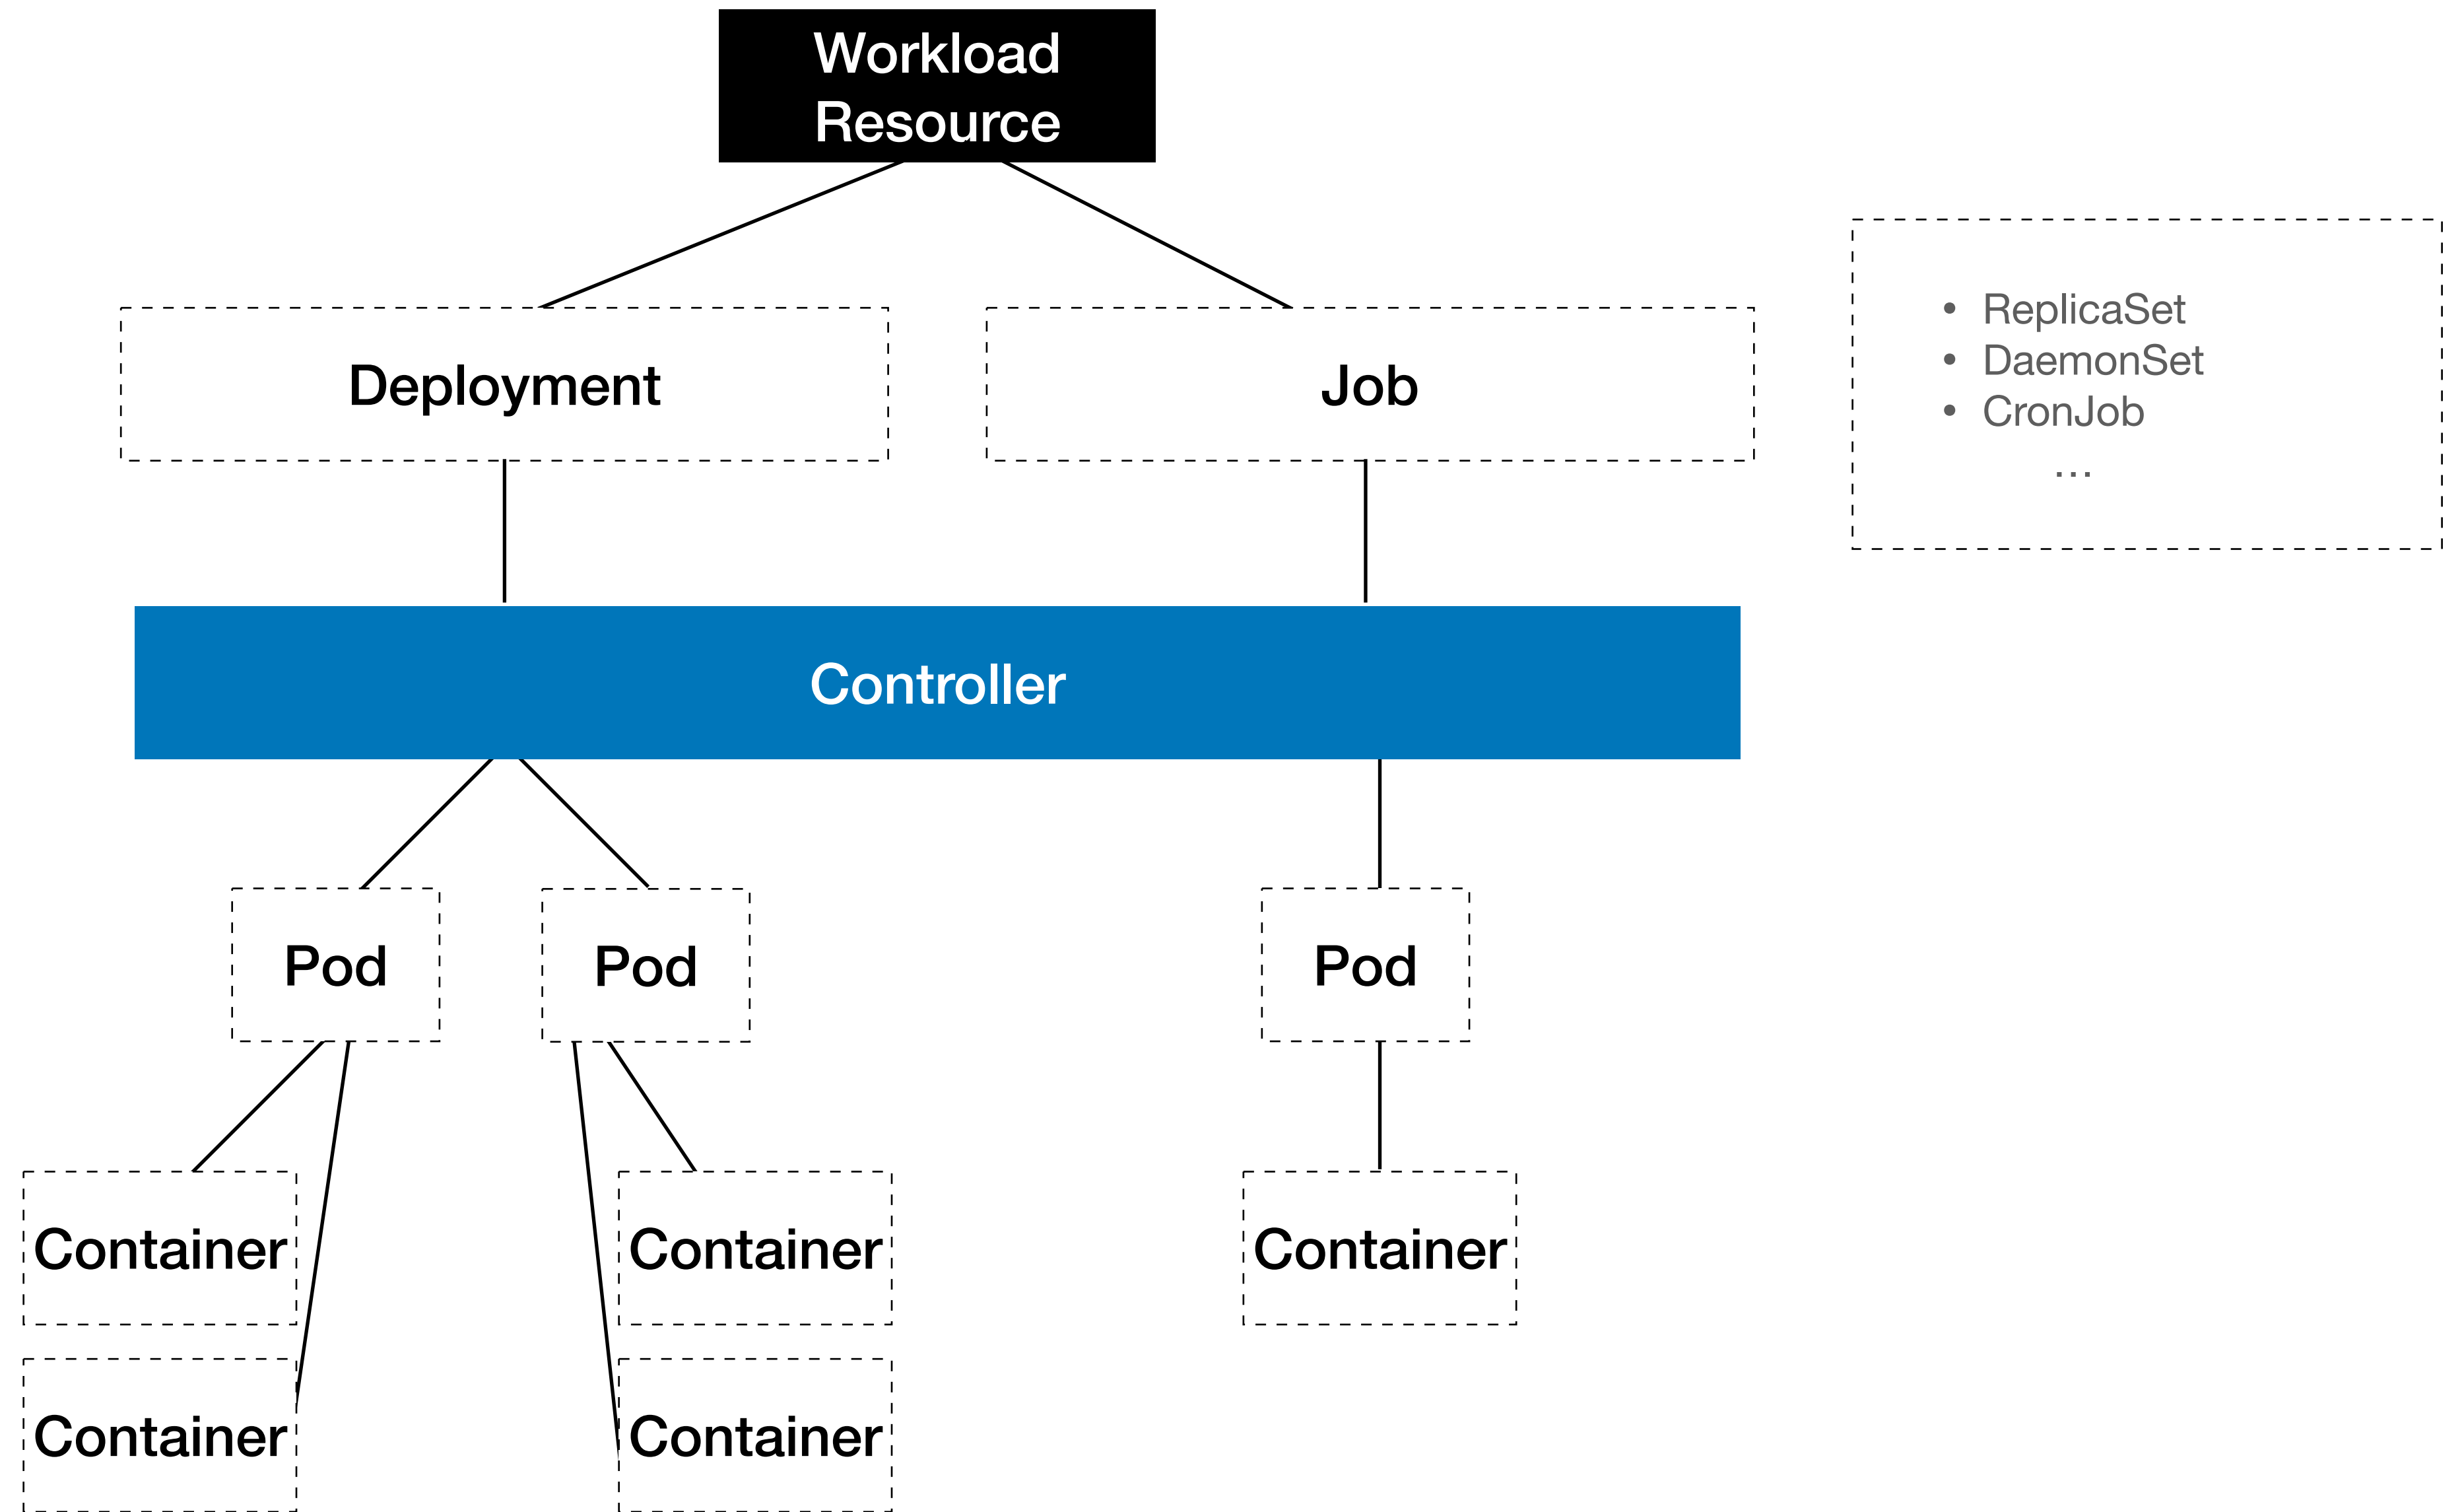
Or using kubectl command to specify the operation (create, replace, etc.), optional flags and at least one file name. The file specified must contain a full definition of the object in YAML or JSON format.

```
kubectl create -f nginx.yaml
```

```
kubectl delete -f nginx.yaml -f redis.yaml
```

<https://kubernetes.io/>

Pod is a smallest deployable compute unit that you can create and manage in Kubernetes



Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
```

```
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
```

```
spec:
  containers:
    - name: nginx
      image: nginx:1.14.2
      ports:
        - containerPort: 80
```

ReplicaSet to be created

Pod template:

Job

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
      - name: pi
        image: perl
        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
      restartPolicy: Never
  backoffLimit: 4
```

Pod template:

Number of retries

CronJob

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
          restartPolicy: OnFailure
```

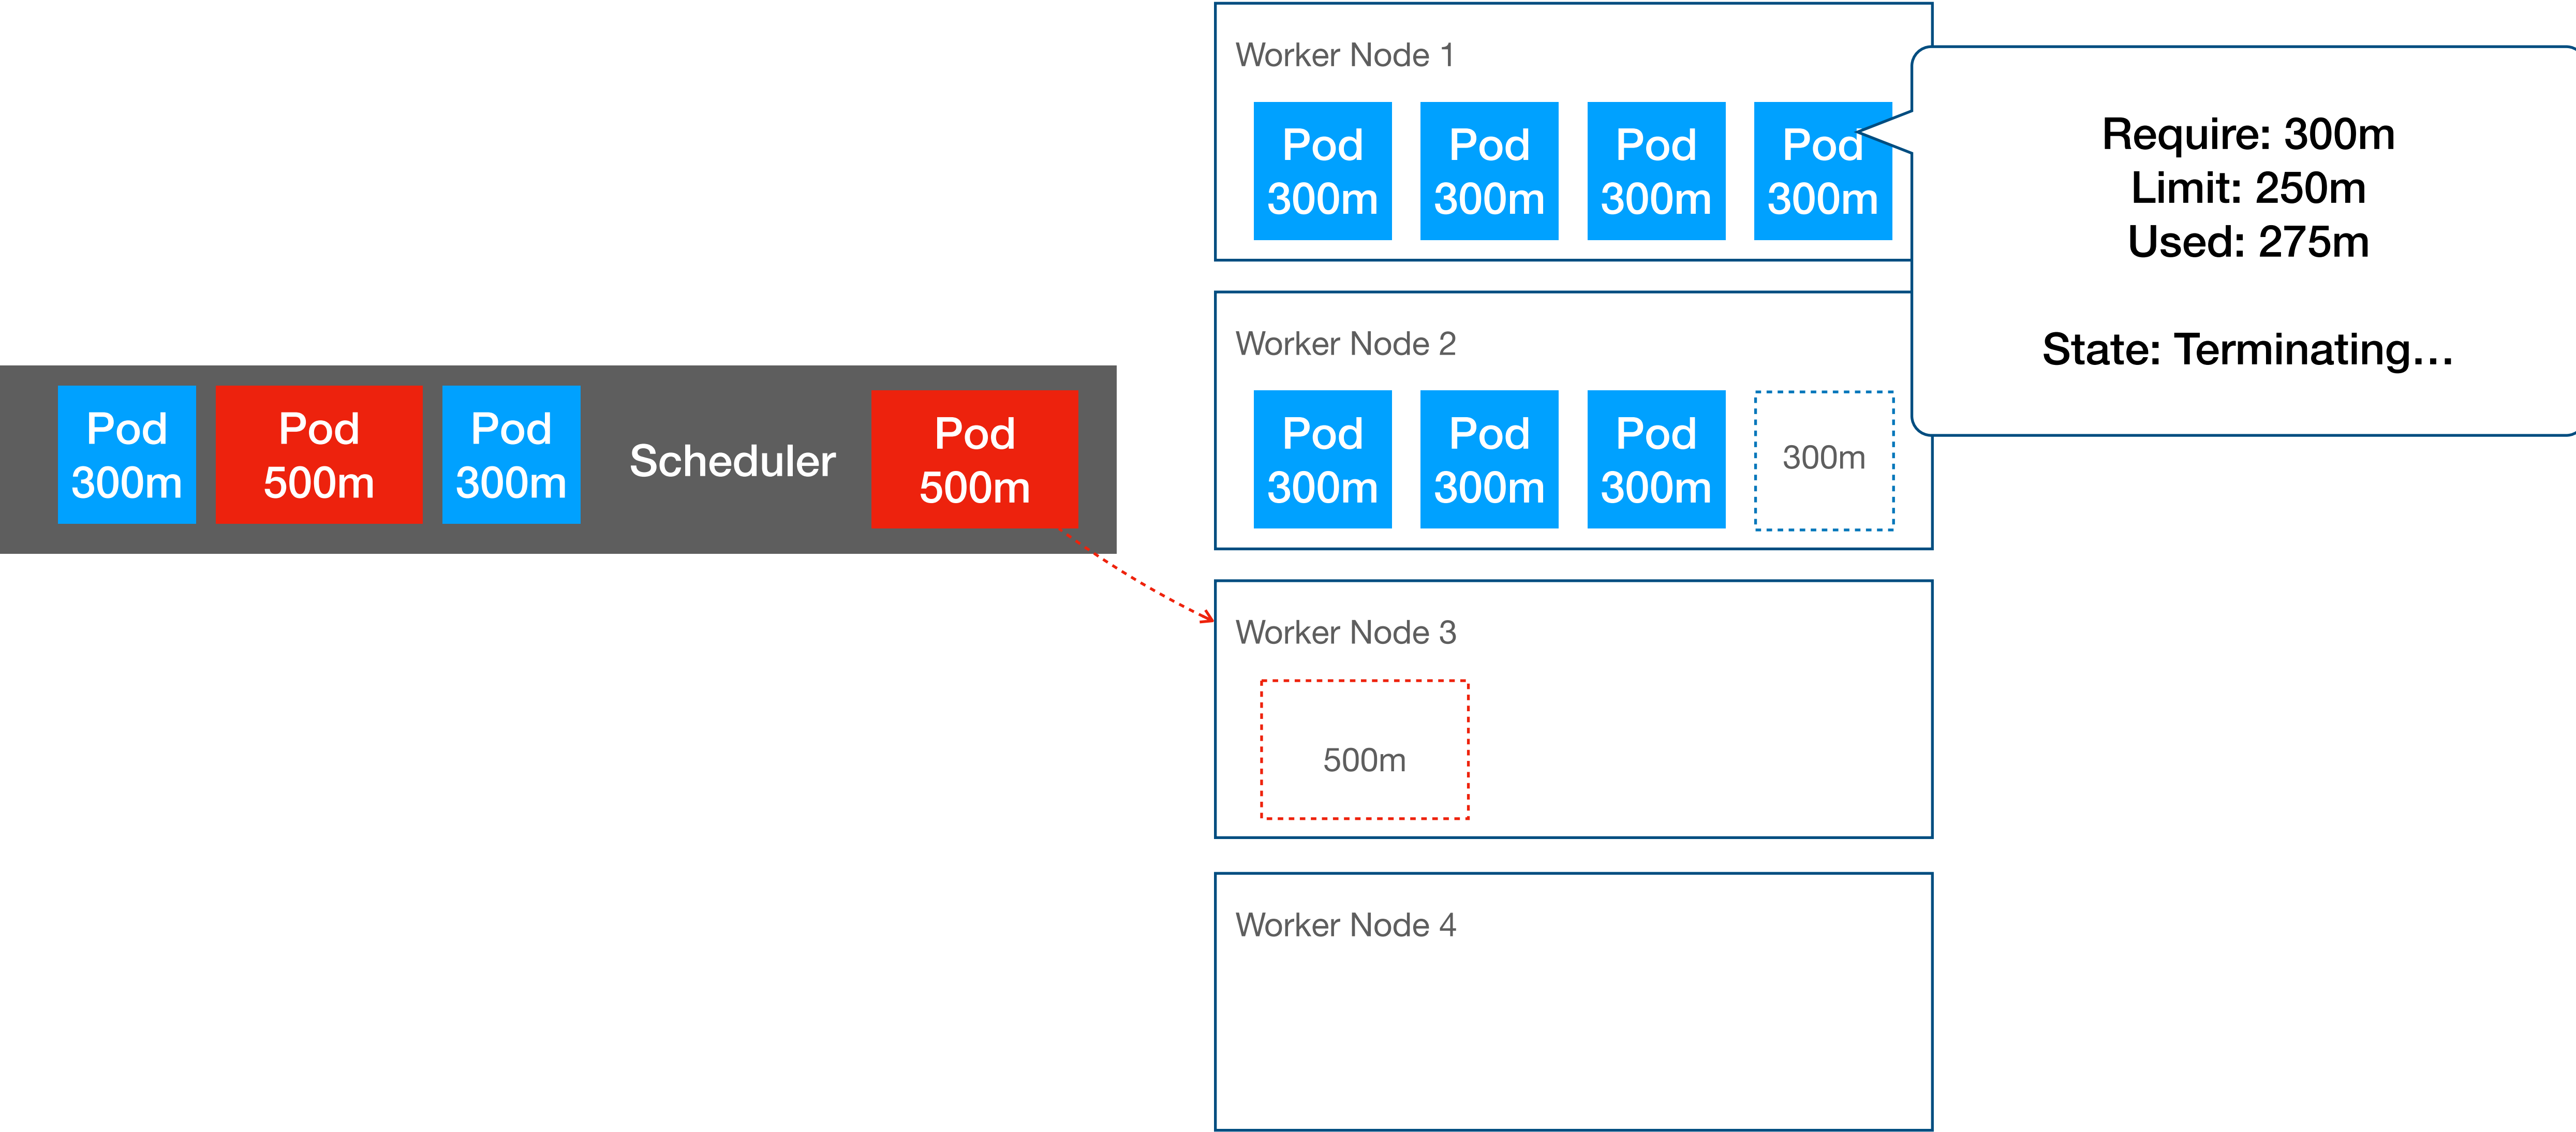
**Use resources requests and limits
available for containers on every Pod right
in your Kubernetes Deployment config**

Resource limits

```
spec:
  containers:
  - name: app
    image: images.my-company.example/app:v4
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
  - name: log-aggregator
    image: images.my-company.example/log-aggregator:v6
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```

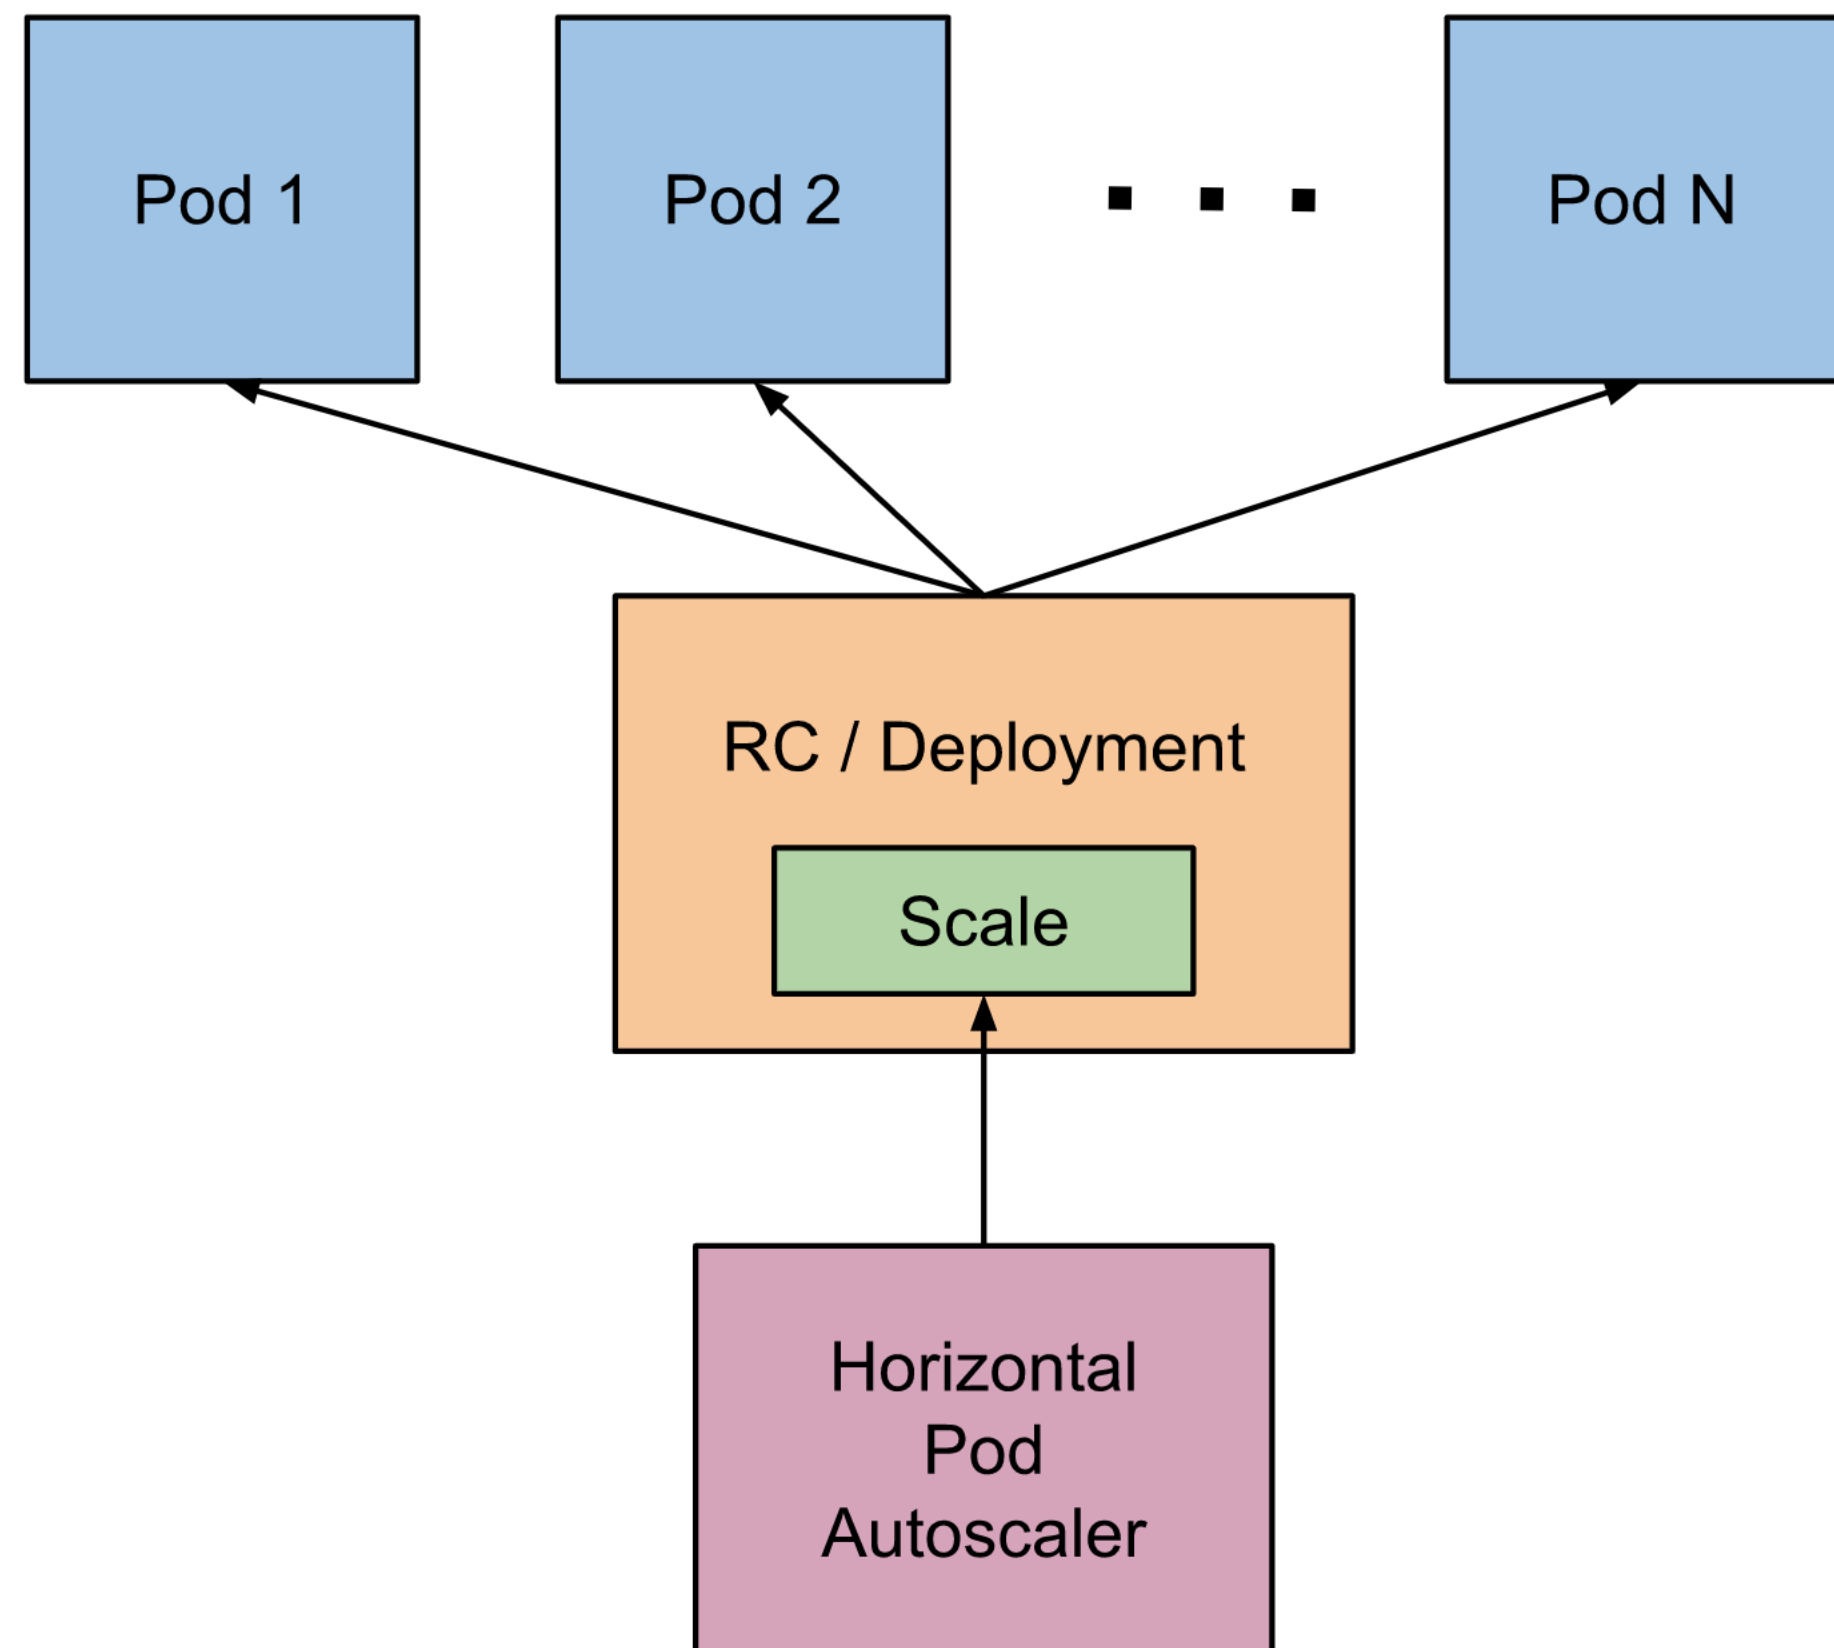
- *request* value is used by scheduler to decide which node to place the Pod on
- *limit* value shows how much resources container is allowed to use
- CPU is specified in units of Kubernetes CPU, where 1 CPU is equivalent of 1 vCPU/core on cloud providers
- 250m (millicpus) = 0.25 CPU

Split between nodes by *require* value



**Resource limitation helps scheduler
to manage containers automatically**

Horizontal Pod Autoscaler automatically scales the number of Pods in your deployment, based on CPU utilisation (or application-provided metrics)

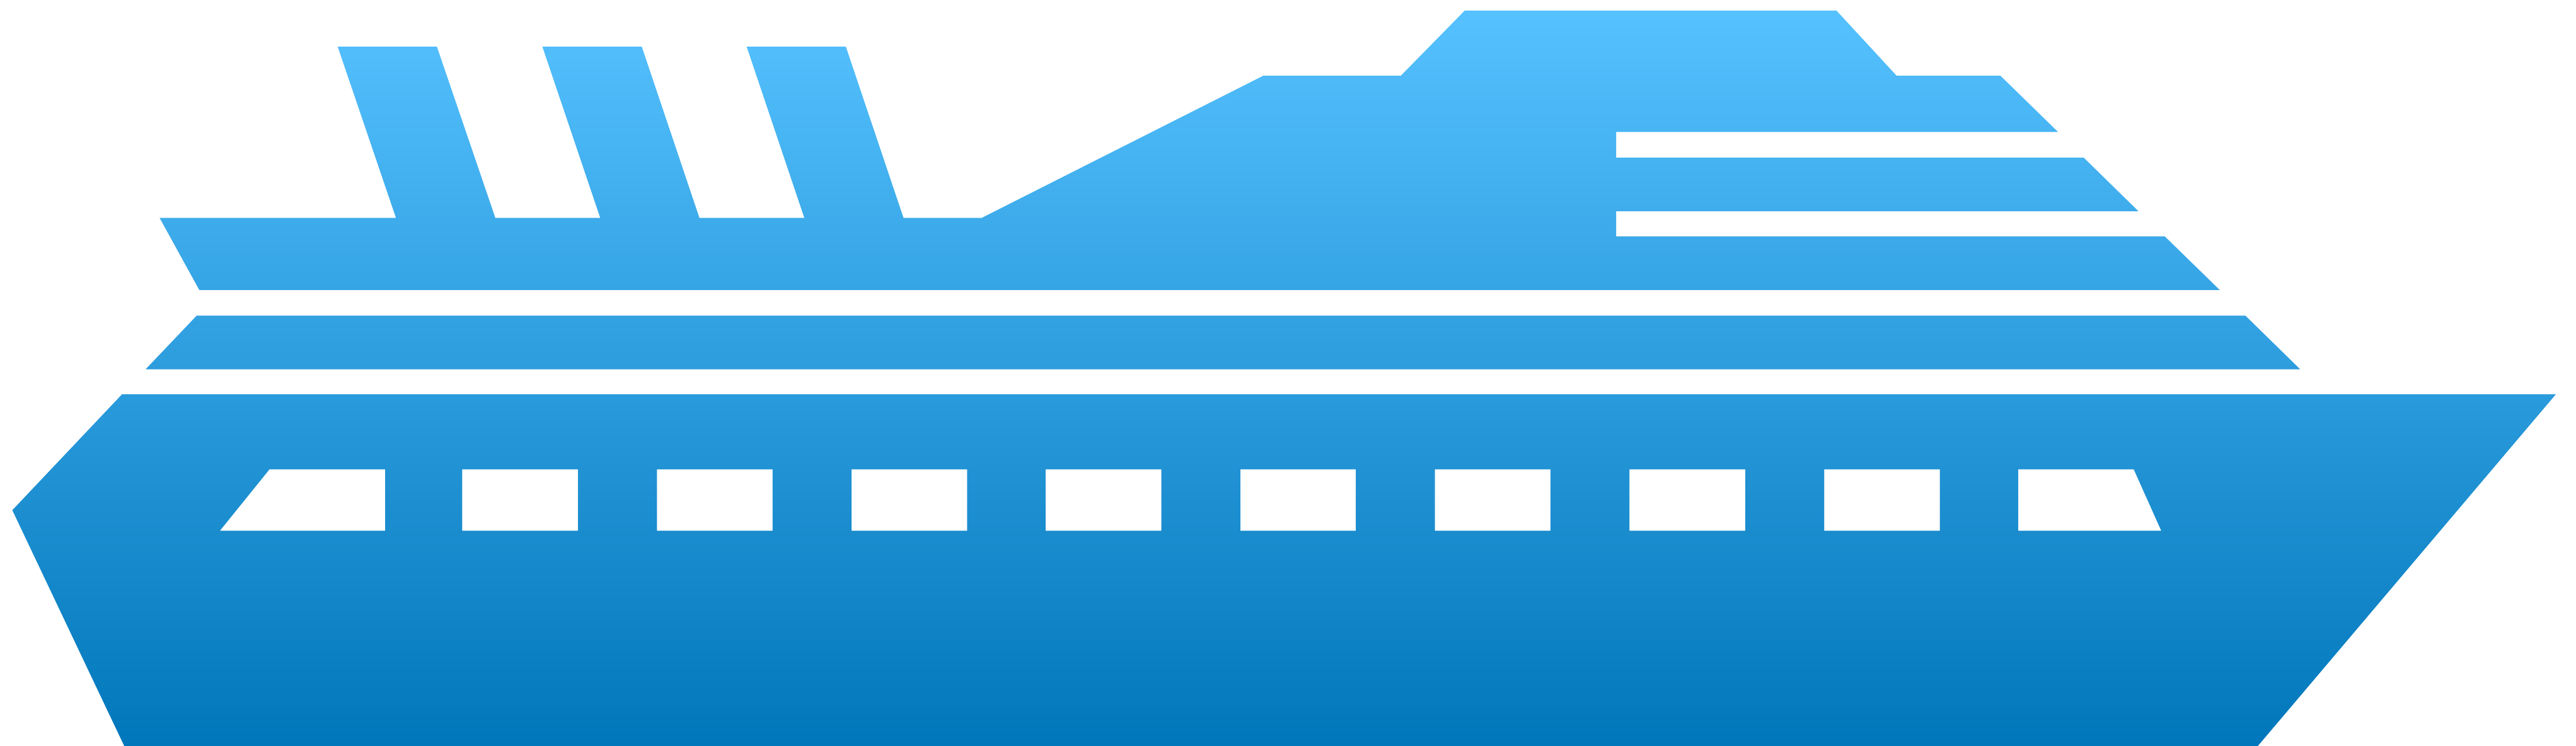


The controller manager periodically queries the resource utilisation and adjusts the number of replicas in a deployment to match the observed average CPU utilisation to the target specified by user.

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

Autoscaler will aim to maintain 50% of average CPU utilisation on all Pods

Demo



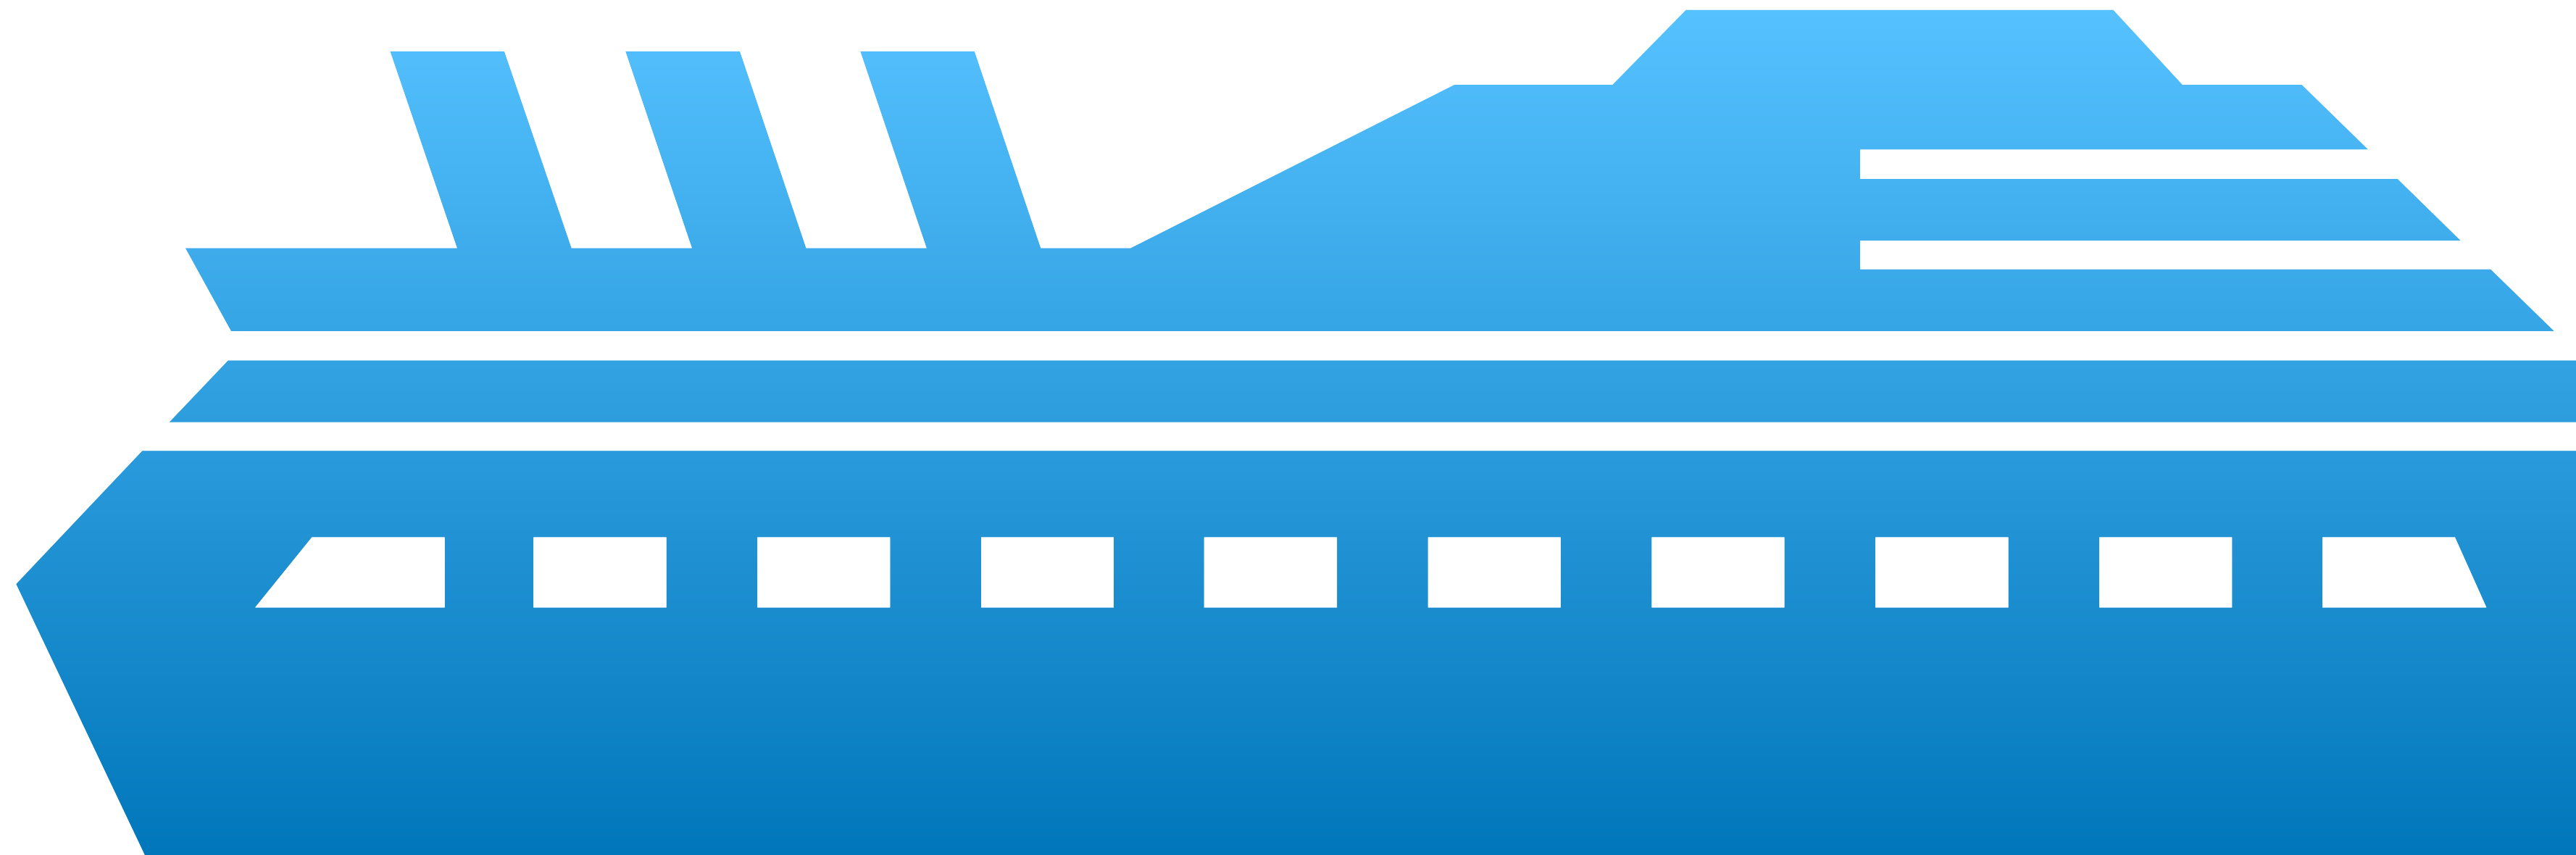


minikube is an easy way to quickly setup Kubernetes cluster locally

```
~$ time minikube start
minikube v1.13.0 on Darwin 10.15.6
Using the docker driver based on user configuration
Starting control plane node minikube in cluster minikube
Creating docker container (CPUs=2, Memory=3892MB) ...
Preparing Kubernetes v1.19.0 on Docker 19.03.8 ...
Verifying Kubernetes components...
Enabled addons: default-storageclass, storage-provisioner
kubectl not found. If you need it, try: 'minikube kubectl -- get pods -A'
Done! kubectl is now configured to use "minikube" by default

Executed in   23.96 secs   fish           external
   usr time    1.66 secs  237.00 micros   1.66 secs
   sys time    0.78 secs   943.00 micros   0.78 secs
```

Questions?



Thank you!

Denys Bulakh

<https://www.linkedin.com/in/denysbulakh/>

